

---

# FontGen

## Generating fonts using a neural net

Matt Gaikema

```
In[1]:= ClearAll["Global`*"]  
SetDirectory[NotebookDirectory[]];
```

I spend a lot of time trying to identify and download fonts, so a program that could do it for me would be wonderful. There are many academic papers on the subject, but none of them have published their code, so I decided to implement it myself and learn more about deep learning along the way.

---

## Data

The data is a whopping 13+ gigabytes and was scraped from the internet by someone else. It can be downloaded [here](#).

```
Import["fonts.hdf5", {"Dimensions"}]
```

```
<| /fonts → {56 443, 62, 64, 64} |>
```

I debugged the neural net using a much smaller dataset, consisting of my personal collection of fonts. The smaller dataset was created by running a script written by the same guy who got the original data.

```
Import["fonts.small.hdf5", {"Dimensions"}]  
f = Import["fonts.small.hdf5", {"Datasets", "fonts"}];  
Image[f[[1]][[1]]]
```

```
<| /fonts → {21, 62, 64, 64} |>
```



---

## Network

Each letter is 64x64 pixels.

```
(* We use 4 letters as input... *)
4 * 64^2
(* ... and get 62 letters as output. *)
62 * 64^2
```

```
16 384
```

```
253 952
```

There are  $4 * 64^2 = 16,384$  input neurons, and  $62 * 64^2 = 253,952$  output neurons. The letters chosen for the inputs are “B”, “A”, “S”, and “Q”.

```
(* http://bit.ly/2oeHyZX *)
Range[0, 25];
letters = AssociationThread[ToUpperCase[#] & /@ Alphabet[], %];
Lookup[letters, {"B", "A", "S", "Q"}]
Image[f[[1]][[#+1]]] & /@ %
```

```
{1, 0, 18, 16}
```

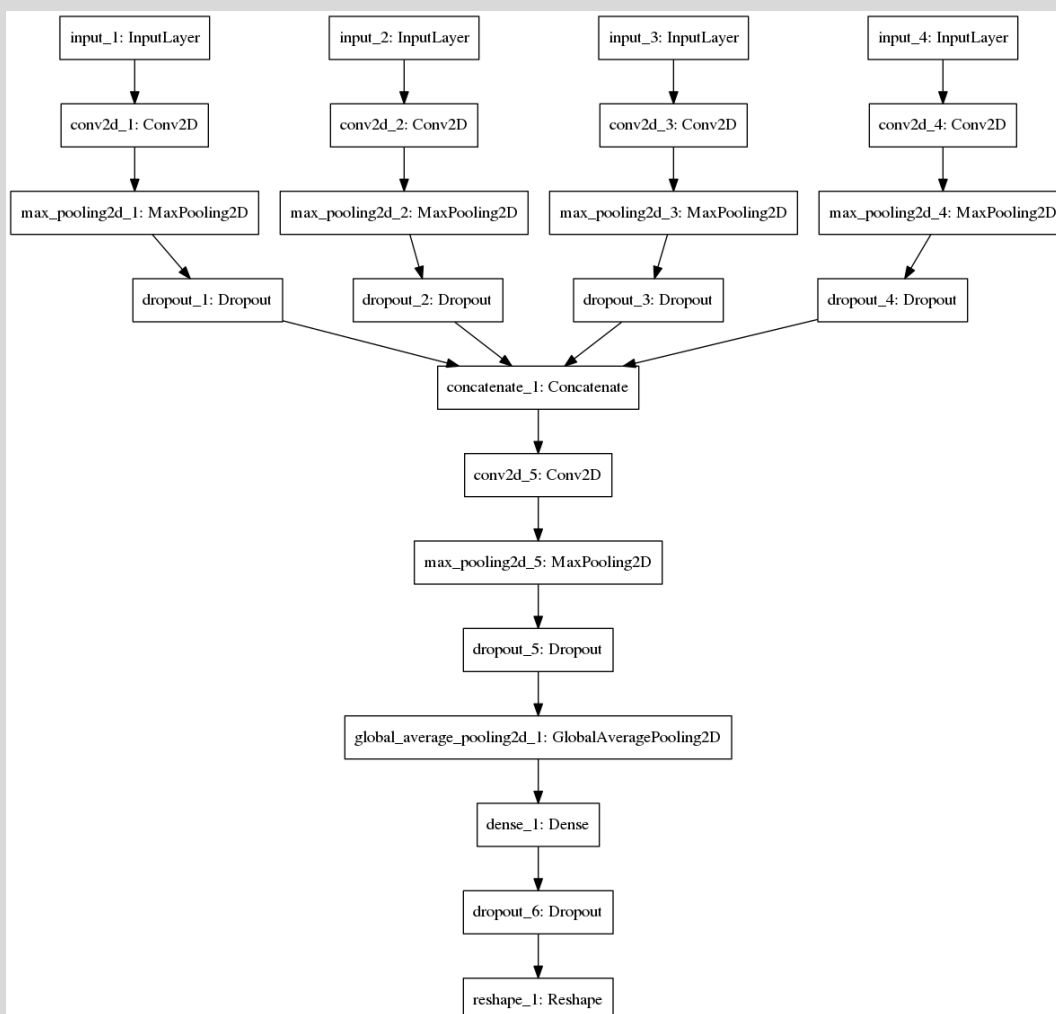


## Architecture

Here is the architecture that I used for the model. Many of the parameters were determined through tedious trial-and-error.

```
(* Created using Keras's `plot_model` function: https://
keras.io/visualization/ *)
Import["img/model.png"]
```

Out[62]=



After training is complete, running `model.save_weights()` saves the model and its weights to an HDF5 file, which can be loaded to recreate the model (see the documentation).

In[63]:=

```
Import["model.hdf5", {"Dimensions"}]
```

Out[63]=

```
<| /conv2d_1/conv2d_1/bias:0 -> {10}, /conv2d_1/conv2d_1/kernel:0 -> {4, 4, 1, 10},
/conv2d_2/conv2d_2/bias:0 -> {10}, /conv2d_2/conv2d_2/kernel:0 -> {4, 4, 1, 10},
/conv2d_3/conv2d_3/bias:0 -> {10}, /conv2d_3/conv2d_3/kernel:0 -> {4, 4, 1, 10},
/conv2d_4/conv2d_4/bias:0 -> {10}, /conv2d_4/conv2d_4/kernel:0 -> {4, 4, 1, 10},
/conv2d_5/conv2d_5/bias:0 -> {8}, /conv2d_5/conv2d_5/kernel:0 -> {8, 8, 40, 8},
/dense_1/dense_1/bias:0 -> {253952}, /dense_1/dense_1/kernel:0 -> {8, 253952} |>
```

## Evaluation


```
In[3]:= (* Transforms a list of image arrays to a list of images. *)
displayFont[f_] := Map[Map[Image, #] &, f];
```

I tested the model on the font Ubuntu Mono, which is available for download [here](#).


```
In[64]:= Import["test.hdf5", {"Dimensions"}]
testInput = Import["test.hdf5", {"input"}];
displayFont[testInput]
output = Import["test.hdf5", {"output"}];
displayFont[output]
```

```
Out[64]= <| /input → {4, 1, 64, 64, 1}, /output → {1, 62, 64, 64} |>
```

```
Out[66]=
```



Out[68]=



You can make out the individual letters, but clearly this isn't an exact recreation. Since I'm still relatively new to deep learning, I'm fairly pleased with these results.

## References

In[59]=

```
(* There is no better way to insert references on Mac/Linux. see http://bit.ly/2lo3Pkc *)
Import["sources.bib"]
```

Out[59]=

```
@inproceedings{paper,
  title = {Learning Typographic Style},
  author = {Shumeet Baluja},
  year = {2016},
  URL = {http://arxiv.org/abs/1603.04000},
  booktitle = {arXiv}
}

@misc{deepfont,
  url =
{https://erikbern.com/2016/01/21/analyzing-50k-fonts-using-deep-neural-networks.html},
  title = {Analyzing 50k fonts using deep neural networks},
  author = {Erik Bernhardsson},
  year = {2016},
  month = JAN
}

@misc{fonsi,
  author = {Fonsi Bonilla},
  howpublished = {Personal Correspondance},
  year = {2018},
  month = DEC
}
```